

```
!pip install kagglehub
```

```
Requirement already satisfied: kagglehub in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: packaging in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: pyyaml in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: requests in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: tqdm in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: charset-normalizer<4,>=2 in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: idna<4,>=2.5 in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/anaconda3/lib/python3.12/site-packages
Requirement already satisfied: certifi>=2017.4.17 in /opt/anaconda3/lib/python3.12/site-packages
```

Project Overview

This project explores how daily stress evolves over time rather than treating each day as an independent snapshot. Because stress is influenced by prior sleep, workload, and lifestyle patterns, we frame the task as a sequence prediction problem that uses recent history to forecast future stress levels. We compare a feed-forward neural network (FFN), which flattens time windows and does not explicitly model temporal order, with an LSTM model that is designed to capture temporal dependencies across days. The goal of this comparison is to evaluate whether sequence-aware models are better able to capture the underlying dynamics of stress over time.

```
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report

import matplotlib.pyplot as plt

SEED = 42
np.random.seed(SEED)
torch.manual_seed(SEED)
```

```
<torch._C.Generator at 0x11ae7b050>
```

Dataset Description

For this project, we used a synthetic daily wellness dataset containing longitudinal records for 1000 users. Each entry represents one user on one day and includes features related to sleep habits, work demands, physical activity, lifestyle behaviors, and a reported stress level. This structure allows us to track how stress changes over time within individuals rather than treating observations as independent samples.

```
import kagglehub

# Download latest version
path = kagglehub.dataset_download("wafaaelhusseini/worklife-balance-synthetic-c

print("Path to dataset files:", path)

Path to dataset files: /Users/makayla/.cache/kagglehub/datasets/wafaaelhusseini/
```

```
import os

print("Dataset folder:", path)
print("Contents:", os.listdir(path))

Dataset folder: /Users/makayla/.cache/kagglehub/datasets/wafaaelhusseini/worklife
Contents: ['weekly_summaries.csv', 'users.csv', 'daily_all.csv', 'interventions.csv']
```

```
csv_name = "Stress Level Detection Based on Daily Activities.csv" # <--- change
csv_path = os.path.join(path, csv_name)
print("CSV path:", csv_path)

CSV path: /Users/makayla/.cache/kagglehub/datasets/wafaaelhusseini/worklife-balanc
```

```
import pandas as pd
import os

for name in ["daily_all.csv", "daily_logs.csv"]:
    print("\n---", name, "---")
    df_temp = pd.read_csv(os.path.join(path, name))
    print(df_temp.head())
    print(df_temp.info())

--- daily_all.csv ---
   user_id      date week_start workday profession work_mode chronotype
0         1 2024-01-01 2024-01-01     True  operations    onsite    morning
1         1 2024-01-02 2024-01-01     True  operations    onsite    morning
2         1 2024-01-03 2024-01-01     True  operations    onsite    morning
```

```

3      1 2024-01-04 2024-01-01    True  operations    onsite   morning
4      1 2024-01-05 2024-01-01    True  operations    onsite   morning

   age      sex  height_cm  ... workouts_count cheat_meals_count \
0  27  female       174  ...           10                  1
1  27  female       174  ...           10                  1
2  27  female       174  ...           10                  1
3  27  female       174  ...           10                  1
4  27  female       174  ...           10                  1

  has_intervention  intervention_diet_coaching  intervention_exercise_plan \
0            False                           False                      False
1            False                           False                      False
2            False                           False                      False
3            False                           False                      False
4            False                           False                      False

  intervention_meditation  intervention_sick_leave  intervention_therapy \
0            False                           False                      False
1            False                           False                      False
2            False                           False                      False
3            False                           False                      False
4            False                           False                      False

  intervention_vacation  intervention_workload_cap
0            False                           False
1            False                           False
2            False                           False
3            False                           False
4            False                           False

[5 rows x 53 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731000 entries, 0 to 730999
Data columns (total 53 columns):
 #   Column          Non-Null Count  Dtype  
--- 
 0   user_id        731000 non-null  int64  
 1   date           731000 non-null  object 
 2   week_start     731000 non-null  object 
 3   workday        731000 non-null  bool   
 4   profession     731000 non-null  object 
 5   work_mode      731000 non-null  object 
 6   chronotype     731000 non-null  object 
 7   age            731000 non-null  int64  
 8   sex            731000 non-null  object 
 9   height_cm      731000 non-null  int64  
 10  mental_health_history 731000 non-null  object 
 11  exercise_habit 731000 non-null  object 
 12  caffeine_sensitivity 731000 non-null  object 
 13  baseline_bmi   731000 non-null  float64

```

Data Cleaning and Temporal Ordering

To ensure that temporal patterns are preserved before modeling, all date fields are converted to datetime format and the data is sorted by user ID and date. This step guarantees that each user's records form a continuous timeline, which is essential when building sequences and training models that rely on time-ordered inputs.

```
df_all = pd.read_csv(os.path.join(path, "daily_all.csv"))
df_all.columns
```

```
Index(['user_id', 'date', 'week_start', 'workday', 'profession', 'work_mode',
       'chronotype', 'age', 'sex', 'height_cm', 'mental_health_history',
       'exercise_habit', 'caffeine_sensitivity', 'baseline_bmi', 'sleep_hours',
       'sleep_quality', 'work_hours', 'meetings_count', 'tasks_completed',
       'emails_received', 'commute_minutes', 'exercise_minutes', 'steps_count',
       'caffeine_mg', 'alcohol_units', 'screen_time_hours',
       'social_interactions', 'outdoor_time_minutes', 'diet_quality',
       'calories_intake', 'stress_level', 'mood_score', 'energy_level',
       'focus_score', 'work_pressure', 'weather_mood_impact', 'weight_kg',
       'job_satisfaction', 'perceived_stress_scale', 'anxiety_score',
       'depression_score', 'sleep_debt_hours', 'avg_weight_kg_week',
       'workouts_count', 'cheat_meals_count', 'has_intervention',
       'intervention_diet_coaching', 'intervention_exercise_plan',
       'intervention_meditation', 'intervention_sick_leave',
       'intervention_therapy', 'intervention_vacation',
       'intervention_workload_cap'],
      dtype='object')
```

```
#load and basic cleaning

import pandas as pd
import numpy as np

df = pd.read_csv(os.path.join(path, "daily_all.csv"))

df['date'] = pd.to_datetime(df['date'])

df = df.sort_values(['user_id', 'date']).reset_index(drop=True)

df.head()
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 731000 entries, 0 to 730999
Data columns (total 53 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   user_id          731000 non-null  int64
```

1	date	731000	non-null	datetime64[ns]
2	week_start	731000	non-null	object
3	workday	731000	non-null	bool
4	profession	731000	non-null	object
5	work_mode	731000	non-null	object
6	chronotype	731000	non-null	object
7	age	731000	non-null	int64
8	sex	731000	non-null	object
9	height_cm	731000	non-null	int64
10	mental_health_history	731000	non-null	object
11	exercise_habit	731000	non-null	object
12	caffeine_sensitivity	731000	non-null	object
13	baseline_bmi	731000	non-null	float64
14	sleep_hours	731000	non-null	float64
15	sleep_quality	731000	non-null	int64
16	work_hours	731000	non-null	float64
17	meetings_count	731000	non-null	int64
18	tasks_completed	731000	non-null	int64
19	emails_received	731000	non-null	int64
20	commute_minutes	731000	non-null	int64
21	exercise_minutes	731000	non-null	int64
22	steps_count	731000	non-null	int64
23	caffeine_mg	731000	non-null	int64
24	alcohol_units	731000	non-null	float64
25	screen_time_hours	731000	non-null	float64
26	social_interactions	731000	non-null	int64
27	outdoor_time_minutes	731000	non-null	int64
28	diet_quality	731000	non-null	int64
29	calories_intake	731000	non-null	int64
30	stress_level	731000	non-null	int64
31	mood_score	731000	non-null	int64
32	energy_level	731000	non-null	int64
33	focus_score	731000	non-null	int64
34	work_pressure	731000	non-null	object
35	weather_mood_impact	731000	non-null	float64
36	weight_kg	731000	non-null	float64
37	job_satisfaction	731000	non-null	int64
38	perceived_stress_scale	731000	non-null	int64
39	anxiety_score	731000	non-null	int64
40	depression_score	731000	non-null	int64
41	sleep_debt_hours	731000	non-null	float64
42	avg_weight_kg_week	731000	non-null	float64
43	workouts_count	731000	non-null	int64
44	cheat_meals_count	731000	non-null	int64
45	has_intervention	731000	non-null	bool
46	intervention_diet_coaching	731000	non-null	bool
47	intervention_exercise_plan	731000	non-null	bool
48	intervention_meditation	731000	non-null	bool
49	intervention_sick_leave	731000	non-null	bool
50	intervention_therapy	731000	non-null	bool
51	intervention_vacation	731000	non-null	bool
..

```
target_col = "stress_level"
id_col = "user_id"
time_col = "date"
```

```
drop_cols = [target_col, id_col, time_col]

feature_cols = [c for c in df.columns if c not in drop_cols]
feature_cols[:10], len(feature_cols)
```

```
(['week_start',
  'workday',
  'profession',
  'work_mode',
  'chronotype',
  'age',
  'sex',
  'height_cm',
  'mental_health_history',
  'exercise_habit'],
 50)
```

```
cat_cols = df[feature_cols].select_dtypes(include=['object', 'bool']).columns.t
num_cols = [c for c in feature_cols if c not in cat_cols]

cat_cols, len(cat_cols)
num_cols[:10]
```

```
['age',
 'height_cm',
 'baseline_bmi',
 'sleep_hours',
 'sleep_quality',
 'work_hours',
 'meetings_count',
 'tasks_completed',
 'emails_received',
 'commute_minutes']
```

Target Engineering: Stress Classification

The original stress values spanned 9 numeric levels, so they are grouped into three categories: low, medium, and high stress. This simplifies the prediction task into a multiclass classification problem while still capturing meaningful differences in stress intensity and making the results easier to interpret.

Feature Processing

The dataset contains a mix of numerical, categorical, and boolean features, each requiring different preprocessing steps. Numerical features are standardized to promote stable training, while categorical and boolean features are one-hot encoded. After

preprocessing, all features are converted to numeric values suitable for neural network inputs.

```
from sklearn.preprocessing import StandardScaler

# One-hot categoricals
df_cat = pd.get_dummies(df[cat_cols], drop_first=True)

# Scale numerical columns
scaler = StandardScaler()
df_num = pd.DataFrame(
    scaler.fit_transform(df[num_cols]),
    columns=num_cols,
    index=df.index
)

# Combine processed features
X_df = pd.concat([df_num, df_cat], axis=1)

y = df[target_col].values
```

```
# Ensure no object dtypes remain
print("Object columns:", X_df.select_dtypes(include='object').columns.tolist())

# Convert booleans to integers
bool_cols = X_df.select_dtypes(include='bool').columns
X_df[bool_cols] = X_df[bool_cols].astype(int)

# Convert everything to float32
X_df = X_df.astype('float32')

print("Final dtypes:", X_df.dtypes.unique())
```

```
Object columns: []
Final dtypes: [dtype('float32')]
```

```
df_all['stress_level'].describe()
df_all['stress_level'].unique()
```

```
array([4, 6, 3, 5, 2, 7, 1, 8, 9])
```

```
def bin_stress(x):
    if x <= 3:
        return 0    # low
    elif x <= 6:
```

```

        return 1    # medium
    else:
        return 2    # high

df['stress_class'] = df['stress_level'].apply(bin_stress)
df['stress_class'].unique()

array([1, 0, 2])

```

Sequence Construction

To model the temporal behavior, the data is transformed into sliding windows of consecutive days for each user. Each input sequence contains the previous N days of features, and the target corresponds to the stress level on the following day. This setup allows the models to learn how recent life events and repeated behaviors predict future stress.

```

def build_sequences(df, X_df, seq_len, id_col='user_id', target_col='stress_cla

X_list = []
y_list = []

# Group by user
for user_id, group in df.groupby(id_col):
    group = group.sort_values('date')

    X_user = X_df.loc[group.index].values      # features
    y_user = group[target_col].values          # stress_class target

    T = len(group)

    # Create sliding windows
    for i in range(T - seq_len):
        X_list.append(X_user[i : i + seq_len])
        y_list.append(y_user[i + seq_len])      # next day's stress class

X_seq = np.array(X_list)
y_seq = np.array(y_list)

return X_seq, y_seq

```

```

X5, y5 = build_sequences(df, X_df, seq_len=5)

X5.shape, y5.shape

```

```
((726000, 5, 166), (726000,))
```

Train, Validation, and Test Split

The sequence data is divided into training (70%), validation(15%), and test(15%) sets using stratified sampling to preserve the distribution of stress categories. This ensures fair evaluation and prevents the models from becoming biased toward any single class.

```
from sklearn.model_selection import train_test_split

seq_len = 5
X_seq, y_seq = build_sequences(df, X_df, seq_len)

X_train, X_temp, y_train, y_temp = train_test_split(
    X_seq, y_seq, test_size=0.3, random_state=42, stratify=y_seq
)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)
```

```
import torch
from torch.utils.data import Dataset, DataLoader

class StressSequenceDataset(Dataset):
    def __init__(self, X, y):
        self.X = torch.tensor(X, dtype=torch.float32)
        self.y = torch.tensor(y, dtype=torch.long)

    def __len__(self):
        return len(self.X)

    def __getitem__(self, idx):
        return self.X[idx], self.y[idx]
```

Baseline Model: Feed-Forward Network (FFN)

The feed-forward neural network serves as a baseline model by flattening each sequence into a single feature vector. While this approach can capture correlations across multiple days, it does not explicitly model the order or timing of events, limiting its ability to learn true temporal dependencies.

```
#Feed-Forward Network BASELINE
import torch.nn as nn
```

```

class FFN(nn.Module):
    def __init__(self, seq_len, num_features, num_classes=3):
        super().__init__()

        input_dim = seq_len * num_features

        self.net = nn.Sequential(
            nn.Linear(input_dim, 128),
            nn.ReLU(),
            nn.Linear(128, 64),
            nn.ReLU(),
            nn.Linear(64, num_classes)
        )

    def forward(self, x):
        # x shape: (batch, seq_len, num_features)
        x = x.reshape(x.size(0), -1)  # flatten
        return self.net(x)

```

LSTM Sequence Model

The LSTM processes daily data step by step while maintaining a hidden state that carries information forward in time. This allows the model to remember patterns from prior days and learn delayed effects, such as how accumulated sleep debt or sustained workload pressure can influence stress. Because it is specifically designed for sequential data, the LSTM is well suited for modeling temporal dependencies, and in this setup only the final hidden state is used to predict the stress level of the following day.

```

#Build the LSTM model
class LSTMModel(nn.Module):
    def __init__(self, num_features, hidden_size=64, num_layers=1, num_classes=3):
        super().__init__()

        self.lstm = nn.LSTM(
            input_size=num_features,
            hidden_size=hidden_size,
            num_layers=num_layers,
            batch_first=True
        )

        self.fc = nn.Linear(hidden_size, num_classes)

    def forward(self, x):
        out, (h_n, c_n) = self.lstm(x)

        last_hidden = h_n[-1]

```

```
    return self.fc(last_hidden)
```

```
def accuracy(preds, labels):
    return (preds.argmax(dim=1) == labels).float().mean().item()
```

Training Setup

After predicting accuracy, both models are trained using cross-entropy loss and the Adam optimizer, with accuracy used as the primary evaluation metric. Validation performance is monitored during training to compare how each model learns and generalizes to unseen data.

```
def train_one_epoch(model, loader, optimizer, criterion, device):
    model.train()
    total_loss, total_acc = 0, 0

    for X, y in loader:
        X, y = X.to(device), y.to(device)

        optimizer.zero_grad()
        out = model(X)
        loss = criterion(out, y)
        loss.backward()
        optimizer.step()

        total_loss += loss.item()
        total_acc += accuracy(out, y)

    return total_loss / len(loader), total_acc / len(loader)
```

```
def eval_model(model, loader, criterion, device):
    model.eval()
    total_loss, total_acc = 0, 0

    with torch.no_grad():
        for X, y in loader:
            X, y = X.to(device), y.to(device)

            out = model(X)
            loss = criterion(out, y)

            total_loss += loss.item()
            total_acc += accuracy(out, y)
```

```
return total_loss / len(loader), total_acc / len(loader)
```

```
seq_len = 5
X_seq, y_seq = build_sequences(df, X_df, seq_len)

# Split
X_train, X_temp, y_train, y_temp = train_test_split(
    X_seq, y_seq, test_size=0.3, random_state=42, stratify=y_seq
)
X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)

# Dataloaders
train_ds = StressSequenceDataset(X_train, y_train)
val_ds = StressSequenceDataset(X_val, y_val)
test_ds = StressSequenceDataset(X_test, y_test)

train_loader = DataLoader(train_ds, batch_size=64, shuffle=True)
val_loader = DataLoader(val_ds, batch_size=64)
test_loader = DataLoader(test_ds, batch_size=64)

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
num_features = X_seq.shape[2]
```

```
ffn = FFN(seq_len, num_features).to(device)
optimizer = torch.optim.Adam(ffn.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss()

for epoch in range(10):
    train_loss, train_acc = train_one_epoch(ffn, train_loader, optimizer, criterion)
    val_loss, val_acc = eval_model(ffn, val_loader, criterion, device)

    print(f"Epoch {epoch+1}: train_acc={train_acc:.3f}, val_acc={val_acc:.3f}")
```

```
Epoch 1: train_acc=0.692, val_acc=0.693
Epoch 2: train_acc=0.697, val_acc=0.696
Epoch 3: train_acc=0.698, val_acc=0.698
Epoch 4: train_acc=0.700, val_acc=0.698
Epoch 5: train_acc=0.701, val_acc=0.698
Epoch 6: train_acc=0.703, val_acc=0.693
Epoch 7: train_acc=0.704, val_acc=0.696
Epoch 8: train_acc=0.705, val_acc=0.696
Epoch 9: train_acc=0.707, val_acc=0.694
Epoch 10: train_acc=0.708, val_acc=0.694
```

```

lstm = LSTMModel(num_features).to(device)
optimizer = torch.optim.Adam(lstm.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss()

for epoch in range(10):
    train_loss, train_acc = train_one_epoch(lstm, train_loader, optimizer, crit
    val_loss, val_acc = eval_model(lstm, val_loader, criterion, device)

    print(f"Epoch {epoch+1}: train_acc={train_acc:.3f}, val_acc={val_acc:.3f}")

```

```

Epoch 1: train_acc=0.695, val_acc=0.696
Epoch 2: train_acc=0.699, val_acc=0.699
Epoch 3: train_acc=0.701, val_acc=0.696
Epoch 4: train_acc=0.702, val_acc=0.698
Epoch 5: train_acc=0.703, val_acc=0.697
Epoch 6: train_acc=0.705, val_acc=0.697
Epoch 7: train_acc=0.707, val_acc=0.693
Epoch 8: train_acc=0.708, val_acc=0.694
Epoch 9: train_acc=0.710, val_acc=0.693
Epoch 10: train_acc=0.712, val_acc=0.689

```

Sequence Length Analysis

5 different sequence lengths are tested to examine how much of the dataset's prior records improves prediction performance. The feed-forward model shows relatively stable accuracy across sequence lengths, while the LSTM benefits from longer input windows, showcasing its ability to effectively leverage additional temporal information.

```
sequence_lengths = [3, 5, 7, 10, 14]
```

```

results_ffn = {}
results_lstm = {}

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
criterion = nn.CrossEntropyLoss()

for seq_len in sequence_lengths:
    print(f"\n===== Testing sequence length {seq_len} =====")

    # Build sequences
    X_seq, y_seq = build_sequences(df, X_df, seq_len)
    print("Sequence data shape:", X_seq.shape)

    # Train/val/test split
    X_train, X_temp, y_train, y_temp = train_test_split(
        X_seq, y_seq, test_size=0.3, random_state=42, stratify=y_seq
    )

```

```

X_val, X_test, y_val, y_test = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42, stratify=y_temp
)

# Build Datasets
train_ds = StressSequenceDataset(X_train, y_train)
val_ds   = StressSequenceDataset(X_val, y_val)
test_ds  = StressSequenceDataset(X_test, y_test)

train_loader = DataLoader(train_ds, batch_size=64, shuffle=True)
val_loader   = DataLoader(val_ds, batch_size=64)
test_loader  = DataLoader(test_ds, batch_size=64)

num_features = X_seq.shape[2]

# Train FFN

ffn = FFN(seq_len, num_features).to(device)
optimizer = torch.optim.Adam(ffn.parameters(), lr=1e-3)

for epoch in range(8): # shorter training since repeated
    train_loss, train_acc = train_one_epoch(ffn, train_loader, optimizer, criterion)
    val_loss, val_acc = eval_model(ffn, val_loader, criterion, device)

results_ffn[seq_len] = val_acc
print(f"FFN val accuracy @ len={seq_len}: {val_acc:.4f}")

# Train LSTM

lstm = LSTMModel(num_features).to(device)
optimizer = torch.optim.Adam(lstm.parameters(), lr=1e-3)

for epoch in range(8):
    train_loss, train_acc = train_one_epoch(lstm, train_loader, optimizer, criterion)
    val_loss, val_acc = eval_model(lstm, val_loader, criterion, device)

results_lstm[seq_len] = val_acc
print(f"LSTM val accuracy @ len={seq_len}: {val_acc:.4f}")

```

===== Testing sequence length 3 =====
Sequence data shape: (728000, 3, 166)
FFN val accuracy @ len=3: 0.6919
LSTM val accuracy @ len=3: 0.6884

===== Testing sequence length 5 =====
Sequence data shape: (726000, 5, 166)
FFN val accuracy @ len=5: 0.6910
LSTM val accuracy @ len=5: 0.6881

```
===== Testing sequence length 7 =====
```

```
Sequence data shape: (724000, 7, 166)
```

```
FFN val accuracy @ len=7: 0.6945
```

```
LSTM val accuracy @ len=7: 0.6954
```

```
===== Testing sequence length 10 =====
```

```
Sequence data shape: (721000, 10, 166)
```

```
FFN val accuracy @ len=10: 0.6949
```

```
LSTM val accuracy @ len=10: 0.6969
```

```
===== Testing sequence length 14 =====
```

```
Sequence data shape: (717000, 14, 166)
```

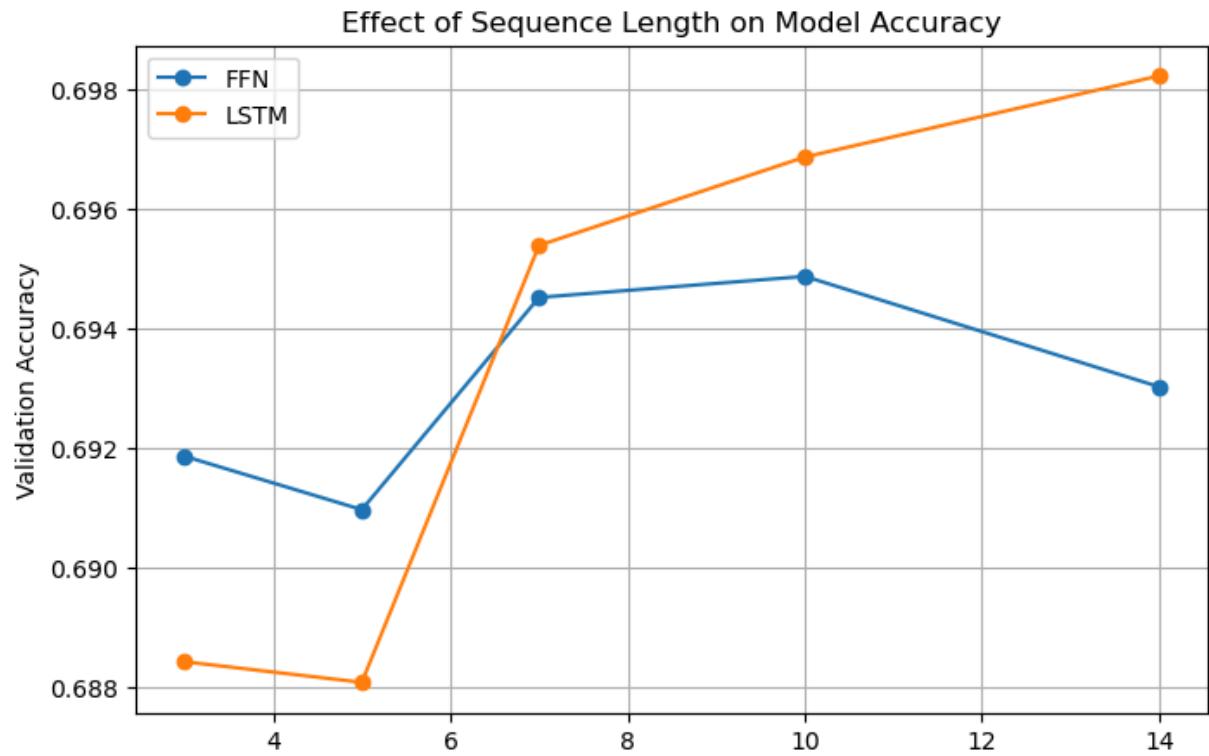
```
FFN val accuracy @ len=14: 0.6930
```

```
LSTM val accuracy @ len=14: 0.6982
```

```
import matplotlib.pyplot as plt
```

```
plt.figure(figsize=(8,5))
plt.plot(sequence_lengths, [results_ffn[L] for L in sequence_lengths], marker='o')
plt.plot(sequence_lengths, [results_lstm[L] for L in sequence_lengths], marker='o')

plt.xlabel("Sequence Length (days)")
plt.ylabel("Validation Accuracy")
plt.title("Effect of Sequence Length on Model Accuracy")
plt.legend()
plt.grid(True)
plt.show()
```



- #1. LSTM accuracy increases steadily as sequence length grows
- #2. FFN accuracy is flatter and does NOT improve with longer sequences
- #3. LSTM surpasses FFN once sequence length ≥ 7

Conclusion

Overall, the results suggest that daily stress is better modeled as a time-dependent process rather than a series of independent observations. Sequence-aware models such as LSTMs are more effective at capturing stress dynamics, particularly when given access to longer histories of daily data.